

CS270: Advanced Operating Systems

Course Project on File System Implementation

Gareth George Thomas Schibler Nazmus Saquib

Graduate Students
Department of Computer Science
University of California Santa Barbara

December 3, 2018



Outline

- 1 Introduction
- 2 Architecture
- 3 Beyond Basics
- 4 Challenges
- 5 Performance Benchmark
- 6 Conclusion
- 7 Questions

Introduction: Design Goals

- Iterative approach to development
- Encapsulation
- Testable
 - object-oriented design
 - unit test for each object
 - integration test

Introduction: Design Goals

- Iterative approach to development
- Encapsulation
- Testable
 - object-oriented design
 - unit test for each object
 - integration test

Introduction: Design Goals

- Iterative approach to development
- Encapsulation
- Testable
 - object-oriented design
 - unit test for each object
 - integration test

Introduction: Design Goals

- Iterative approach to development
- Encapsulation
- Testable
 - object-oriented design
 - unit test for each object
 - integration test

Introduction: Design Goals

- Iterative approach to development
- Encapsulation
- Testable
 - object-oriented design
 - unit test for each object
 - integration test

Introduction: Design Goals

- Iterative approach to development
- Encapsulation
- Testable
 - object-oriented design
 - unit test for each object
 - integration test

Introduction: Design Goals (*Cntd.*)

- General data structures

- superblock
- inode table
- segment manager (LFS)

- Memory mapped files

- performance gain at the cost of reliability
- acceptable for high performance system
- can be seen in Mach¹

¹Accetta et al., "Mach: A new kernel foundation for UNIX development"

Introduction: Design Goals (*Cntd.*)

- General data structures
 - superblock
 - inode table
 - segment manager (LFS)
- Memory mapped files
 - performance gain at the cost of reliability
 - acceptable for high performance system
 - can be seen in Mach¹

¹Accetta et al., "Mach: A new kernel foundation for UNIX development"

Introduction: Design Goals (*Cntd.*)

- General data structures
 - superblock
 - inode table
 - segment manager (LFS)
- Memory mapped files
 - performance gain at the cost of reliability
 - acceptable for high performance system
 - can be seen in Mach¹

¹Accetta et al., "Mach: A new kernel foundation for UNIX development"

Introduction: Design Goals (*Cntd.*)

- General data structures
 - superblock
 - inode table
 - segment manager (LFS)
- Memory mapped files
 - performance gain at the cost of reliability
 - acceptable for high performance system
 - can be seen in Mach¹

¹Accetta et al., "Mach: A new kernel foundation for UNIX development"

Introduction: Design Goals (*Cntd.*)

- General data structures
 - superblock
 - inode table
 - segment manager (LFS)
- Memory mapped files
 - performance gain at the cost of reliability
 - acceptable for high performance system
 - can be seen in Mach¹

¹Accetta et al., "Mach: A new kernel foundation for UNIX development" 

Introduction: Design Goals (*Cntd.*)

- General data structures
 - superblock
 - inode table
 - segment manager (LFS)
- Memory mapped files
 - performance gain at the cost of reliability
 - acceptable for high performance system
 - can be seen in Mach¹

¹Accetta et al., "Mach: A new kernel foundation for UNIX development" 

Introduction: Design Goals (*Cntd.*)

- General data structures
 - superblock
 - inode table
 - segment manager (LFS)
- Memory mapped files
 - performance gain at the cost of reliability
 - acceptable for high performance system
 - can be seen in Mach¹

¹Accetta et al., "Mach: A new kernel foundation for UNIX development" 

Introduction: Design Goals (*Cntd.*)

- General data structures
 - superblock
 - inode table
 - segment manager (LFS)
- Memory mapped files
 - performance gain at the cost of reliability
 - acceptable for high performance system
 - can be seen in Mach¹

¹Accetta et al., “Mach: A new kernel foundation for UNIX development”

Introduction: Memory Mapped Files

- + Efficient paging as kernel handles it
- Lesser direct control
- Restricts fine-grained control over writes

Introduction: Memory Mapped Files

- + Efficient paging as kernel handles it
- Lesser direct control
- Restricts fine-grained control over writes

Introduction: Memory Mapped Files

- + Efficient paging as kernel handles it
- Lesser direct control
- Restricts fine-grained control over writes

Outline

- 1 Introduction
- 2 Architecture**
- 3 Beyond Basics
- 4 Challenges
- 5 Performance Benchmark
- 6 Conclusion
- 7 Questions

Architecture: Hybrid Log Structured File System (LFS)

- Superblock

- holds references to all other file system data structures
- initializes
 - the inode table
 - the segment controller

- Inodes

- closely resembles inodes from UNIX FFS²
- each inode contains 8 direct, 1 indirect, 1 double indirect, and 1 triple indirect blocks
- additionally complicated by the need for an algorithm that allows moving both data blocks and indirect pages belonging to an inode

²McKusick et al., "A fast file system for UNIX".

Architecture: Hybrid Log Structured File System (LFS)

- Superblock

- holds references to all other file system data structures
- initializes
 - the inode table
 - the segment controller

- Inodes

- closely resembles inodes from UNIX FFS²
- each inode contains 8 direct, 1 indirect, 1 double indirect, and 1 triple indirect blocks
- additionally complicated by the need for an algorithm that allows moving both data blocks and indirect pages belonging to an inode

²McKusick et al., "A fast file system for UNIX".

Architecture: Hybrid Log Structured File System (LFS)

- Superblock

- holds references to all other file system data structures
- initializes
 - the inode table
 - the segment controller

- Inodes

- closely resembles inodes from UNIX FFS²
- each inode contains 8 direct, 1 indirect, 1 double indirect, and 1 triple indirect blocks
- additionally complicated by the need for an algorithm that allows moving both data blocks and indirect pages belonging to an inode

²McKusick et al., "A fast file system for UNIX".

Architecture: Hybrid Log Structured File System (LFS)

- Superblock

- holds references to all other file system data structures
- initializes
 - the inode table
 - the segment controller

- Inodes

- closely resembles inodes from UNIX FFS²
- each inode contains 8 direct, 1 indirect, 1 double indirect, and 1 triple indirect blocks
- additionally complicated by the need for an algorithm that allows moving both data blocks and indirect pages belonging to an inode

²McKusick et al., "A fast file system for UNIX".

Architecture: Hybrid Log Structured File System (LFS)

- Superblock

- holds references to all other file system data structures
- initializes
 - the inode table
 - the segment controller

- Inodes

- closely resembles inodes from UNIX FFS²
- each inode contains 8 direct, 1 indirect, 1 double indirect, and 1 triple indirect blocks
- additionally complicated by the need for an algorithm that allows moving both data blocks and indirect pages belonging to an inode

²McKusick et al., "A fast file system for UNIX".

Architecture: Hybrid Log Structured File System (LFS)

- Superblock

- holds references to all other file system data structures
- initializes
 - the inode table
 - the segment controller

- Inodes

- closely resembles inodes from UNIX FFS²
- each inode contains 8 direct, 1 indirect, 1 double indirect, and 1 triple indirect blocks
- additionally complicated by the need for an algorithm that allows moving both data blocks and indirect pages belonging to an inode

²McKusick et al., “A fast file system for UNIX”.

Architecture: Hybrid Log Structured File System (LFS)

- Superblock

- holds references to all other file system data structures
- initializes
 - the inode table
 - the segment controller

- Inodes

- closely resembles inodes from UNIX FFS²
- each inode contains 8 direct, 1 indirect, 1 double indirect, and 1 triple indirect blocks
- additionally complicated by the need for an algorithm that allows moving both data blocks and indirect pages belonging to an inode

²McKusick et al., “A fast file system for UNIX”.

Architecture: Hybrid Log Structured File System (LFS)

- Superblock

- holds references to all other file system data structures
- initializes
 - the inode table
 - the segment controller

- Inodes

- closely resembles inodes from UNIX FFS²
- each inode contains 8 direct, 1 indirect, 1 double indirect, and 1 triple indirect blocks
- additionally complicated by the need for an algorithm that allows moving both data blocks and indirect pages belonging to an inode

²McKusick et al., “A fast file system for UNIX”.

Architecture: Hybrid Log Structured File System (LFS)

- Superblock

- holds references to all other file system data structures
- initializes
 - the inode table
 - the segment controller

- Inodes

- closely resembles inodes from UNIX FFS²
- each inode contains 8 direct, 1 indirect, 1 double indirect, and 1 triple indirect blocks
- additionally complicated by the need for an algorithm that allows moving both data blocks and indirect pages belonging to an inode

²McKusick et al., “A fast file system for UNIX”.

Outline

- 1 Introduction
- 2 Architecture
- 3 Beyond Basics**
- 4 Challenges
- 5 Performance Benchmark
- 6 Conclusion
- 7 Questions

- Segment controller

- heavy lifting data structure
- manages a pool of segments
- tracks segment utilization for efficient GC

- Segment

- each segment maps its data blocks to the inodes that own them
- keeps a write head that allows for constant time chunk allocation
- finding and garbage collecting low utilization chunks can be costly, however

- Experiences with LFS

- + LFS dramatically decreased time to allocate chunks
- + Faster sequential writes due to constant time chunk allocation
- Slower updates due to poor utilization of page cache when doing copy on write for updates
- Introduces costly GC operation

Beyond Basics

- Segment controller
 - heavy lifting data structure
 - manages a pool of segments
 - tracks segment utilization for efficient GC
- Segment
 - each segment maps its data blocks to the inodes that own them
 - keeps a write head that allows for constant time chunk allocation
 - finding and garbage collecting low utilization chunks can be costly, however
- Experiences with LFS
 - + LFS dramatically decreased time to allocate chunks
 - + Faster sequential writes due to constant time chunk allocation
 - Slower updates due to poor utilization of page cache when doing copy on write for updates
 - Introduces costly GC operation

Beyond Basics

- Segment controller
 - heavy lifting data structure
 - manages a pool of segments
 - tracks segment utilization for efficient GC
- Segment
 - each segment maps its data blocks to the inodes that own them
 - keeps a write head that allows for constant time chunk allocation
 - finding and garbage collecting low utilization chunks can be costly, however
- Experiences with LFS
 - + LFS dramatically decreased time to allocate chunks
 - + Faster sequential writes due to constant time chunk allocation
 - Slower updates due to poor utilization of page cache when doing copy on write for updates
 - Introduces costly GC operation

Beyond Basics

- Segment controller
 - heavy lifting data structure
 - manages a pool of segments
 - tracks segment utilization for efficient GC
- Segment
 - each segment maps its data blocks to the inodes that own them
 - keeps a write head that allows for constant time chunk allocation
 - finding and garbage collecting low utilization chunks can be costly, however
- Experiences with LFS
 - + LFS dramatically decreased time to allocate chunks
 - + Faster sequential writes due to constant time chunk allocation
 - Slower updates due to poor utilization of page cache when doing copy on write for updates
 - Introduces costly GC operation

Beyond Basics

- Segment controller
 - heavy lifting data structure
 - manages a pool of segments
 - tracks segment utilization for efficient GC
- Segment
 - each segment maps its data blocks to the inodes that own them
 - keeps a write head that allows for constant time chunk allocation
 - finding and garbage collecting low utilization chunks can be costly, however
- Experiences with LFS
 - + LFS dramatically decreased time to allocate chunks
 - + Faster sequential writes due to constant time chunk allocation
 - Slower updates due to poor utilization of page cache when doing copy on write for updates
 - Introduces costly GC operation

Beyond Basics

- Segment controller
 - heavy lifting data structure
 - manages a pool of segments
 - tracks segment utilization for efficient GC
- Segment
 - each segment maps its data blocks to the inodes that own them
 - keeps a write head that allows for constant time chunk allocation
 - finding and garbage collecting low utilization chunks can be costly, however
- Experiences with LFS
 - + LFS dramatically decreased time to allocate chunks
 - + Faster sequential writes due to constant time chunk allocation
 - Slower updates due to poor utilization of page cache when doing copy on write for updates
 - Introduces costly GC operation

Beyond Basics

- Segment controller
 - heavy lifting data structure
 - manages a pool of segments
 - tracks segment utilization for efficient GC
- Segment
 - each segment maps its data blocks to the inodes that own them
 - keeps a write head that allows for constant time chunk allocation
 - finding and garbage collecting low utilization chunks can be costly, however
- Experiences with LFS
 - + LFS dramatically decreased time to allocate chunks
 - + Faster sequential writes due to constant time chunk allocation
 - Slower updates due to poor utilization of page cache when doing copy on write for updates
 - Introduces costly GC operation

Beyond Basics

- Segment controller
 - heavy lifting data structure
 - manages a pool of segments
 - tracks segment utilization for efficient GC
- Segment
 - each segment maps its data blocks to the inodes that own them
 - keeps a write head that allows for constant time chunk allocation
 - finding and garbage collecting low utilization chunks can be costly, however
- Experiences with LFS
 - + LFS dramatically decreased time to allocate chunks
 - + Faster sequential writes due to constant time chunk allocation
 - Slower updates due to poor utilization of page cache when doing copy on write for updates
 - Introduces costly GC operation

Beyond Basics

- Segment controller
 - heavy lifting data structure
 - manages a pool of segments
 - tracks segment utilization for efficient GC
- Segment
 - each segment maps its data blocks to the inodes that own them
 - keeps a write head that allows for constant time chunk allocation
 - finding and garbage collecting low utilization chunks can be costly, however
- Experiences with LFS
 - + LFS dramatically decreased time to allocate chunks
 - + Faster sequential writes due to constant time chunk allocation
 - Slower updates due to poor utilization of page cache when doing copy on write for updates
 - Introduces costly GC operation

Beyond Basics

- Segment controller
 - heavy lifting data structure
 - manages a pool of segments
 - tracks segment utilization for efficient GC
- Segment
 - each segment maps its data blocks to the inodes that own them
 - keeps a write head that allows for constant time chunk allocation
 - finding and garbage collecting low utilization chunks can be costly, however
- Experiences with LFS
 - + LFS dramatically decreased time to allocate chunks
 - + Faster sequential writes due to constant time chunk allocation
 - Slower updates due to poor utilization of page cache when doing copy on write for updates
 - Introduces costly GC operation

Beyond Basics

- Segment controller
 - heavy lifting data structure
 - manages a pool of segments
 - tracks segment utilization for efficient GC
- Segment
 - each segment maps its data blocks to the inodes that own them
 - keeps a write head that allows for constant time chunk allocation
 - finding and garbage collecting low utilization chunks can be costly, however
- Experiences with LFS
 - + LFS dramatically decreased time to allocate chunks
 - + Faster sequential writes due to constant time chunk allocation
 - Slower updates due to poor utilization of page cache when doing copy on write for updates
 - Introduces costly GC operation

Beyond Basics

- Segment controller
 - heavy lifting data structure
 - manages a pool of segments
 - tracks segment utilization for efficient GC
- Segment
 - each segment maps its data blocks to the inodes that own them
 - keeps a write head that allows for constant time chunk allocation
 - finding and garbage collecting low utilization chunks can be costly, however
- Experiences with LFS
 - + LFS dramatically decreased time to allocate chunks
 - + Faster sequential writes due to constant time chunk allocation
 - Slower updates due to poor utilization of page cache when doing copy on write for updates
 - Introduces costly GC operation

Beyond Basics

- Segment controller
 - heavy lifting data structure
 - manages a pool of segments
 - tracks segment utilization for efficient GC
- Segment
 - each segment maps its data blocks to the inodes that own them
 - keeps a write head that allows for constant time chunk allocation
 - finding and garbage collecting low utilization chunks can be costly, however
- Experiences with LFS
 - + LFS dramatically decreased time to allocate chunks
 - + Faster sequential writes due to constant time chunk allocation
 - Slower updates due to poor utilization of page cache when doing copy on write for updates
 - Introduces costly GC operation

Outline

- 1 Introduction
- 2 Architecture
- 3 Beyond Basics
- 4 Challenges**
- 5 Performance Benchmark
- 6 Conclusion
- 7 Questions

Challenges

- Memory mapped interface
- Dropped pointer
- Off by one

Language choice: C++ 11

- + proper exception handling
- + proper memory management

Challenges

- Memory mapped interface
- Dropped pointer
- Off by one

Language choice: C++ 11

- + proper exception handling
- + proper memory management

Challenges

- Memory mapped interface
- Dropped pointer
- Off by one

Language choice: C++ 11

- + proper exception handling
- + proper memory management

Challenges

- Memory mapped interface
- Dropped pointer
- Off by one

Language choice: C++ 11

- + proper exception handling
- + proper memory management

Challenges

- Memory mapped interface
- Dropped pointer
- Off by one


Language choice: C++ 11

- + proper exception handling
- + proper memory management

Outline

- 1 Introduction
- 2 Architecture
- 3 Beyond Basics
- 4 Challenges
- 5 Performance Benchmark**
- 6 Conclusion
- 7 Questions

Performance Benchmark



Outline

- 1 Introduction
- 2 Architecture
- 3 Beyond Basics
- 4 Challenges
- 5 Performance Benchmark
- 6 Conclusion**
- 7 Questions

Conclusion

Outline

- 1 Introduction
- 2 Architecture
- 3 Beyond Basics
- 4 Challenges
- 5 Performance Benchmark
- 6 Conclusion
- 7 Questions**

Questions?