

# CS270: Advanced Operating Systems

## Course Project on File System Implementation

Gareth George   Thomas Schibler   Nazmus Saquib

Graduate Students  
Department of Computer Science  
University of California Santa Barbara

December 2, 2018



# Outline

- 1 Introduction
- 2 Architecture
- 3 Beyond Basics
- 4 Challenges
- 5 Performance Benchmark
- 6 Conclusion
- 7 Questions

# Introduction: Design Goals

- Iterative approach to development
- Encapsulation
- Testable
  - object-oriented design
  - unit test for each object
  - integration test

# Introduction: Design Goals

- Iterative approach to development
- Encapsulation
- Testable
  - object-oriented design
  - unit test for each object
  - integration test

# Introduction: Design Goals

- Iterative approach to development
- Encapsulation
- Testable
  - object-oriented design
  - unit test for each object
  - integration test

# Introduction: Design Goals

- Iterative approach to development
- Encapsulation
- Testable
  - object-oriented design
  - unit test for each object
  - integration test

# Introduction: Design Goals

- Iterative approach to development
- Encapsulation
- Testable
  - object-oriented design
  - unit test for each object
  - integration test

# Introduction: Design Goals

- Iterative approach to development
- Encapsulation
- Testable
  - object-oriented design
  - unit test for each object
  - integration test



# Introduction: Design Goals (*Cntd.*)

- General data structures

- superblock
- inode table
- segment manager (LFS)

- Memory mapped files

- performance gain at the cost of reliability
- acceptable for high performance system
- can be seen in Mach

# Introduction: Design Goals (*Cntd.*)

- General data structures
  - superblock
  - inode table
  - segment manager (LFS)
- Memory mapped files
  - performance gain at the cost of reliability
  - acceptable for high performance system
  - can be seen in Mach

# Introduction: Design Goals (*Cntd.*)

- General data structures
  - superblock
  - inode table
  - segment manager (LFS)
- Memory mapped files
  - performance gain at the cost of reliability
  - acceptable for high performance system
  - can be seen in Mach

# Introduction: Design Goals (*Cntd.*)

- General data structures
  - superblock
  - inode table
  - segment manager (LFS)
- Memory mapped files
  - performance gain at the cost of reliability
  - acceptable for high performance system
  - can be seen in Mach

# Introduction: Design Goals (*Cntd.*)

- General data structures
  - superblock
  - inode table
  - segment manager (LFS)
- Memory mapped files
  - performance gain at the cost of reliability
  - acceptable for high performance system
  - can be seen in Mach

# Introduction: Design Goals (*Cntd.*)

- General data structures
  - superblock
  - inode table
  - segment manager (LFS)
- Memory mapped files
  - performance gain at the cost of reliability
  - acceptable for high performance system
  - can be seen in Mach

# Introduction: Design Goals (*Cntd.*)

- General data structures
  - superblock
  - inode table
  - segment manager (LFS)
- Memory mapped files
  - performance gain at the cost of reliability
  - acceptable for high performance system
  - can be seen in Mach

# Introduction: Design Goals (*Cntd.*)

- General data structures
  - superblock
  - inode table
  - segment manager (LFS)
- Memory mapped files
  - performance gain at the cost of reliability
  - acceptable for high performance system
  - can be seen in Mach



# Introduction: Memory Mapped Files

- + Efficient paging as kernel handles it
- Lesser direct control
- Restricts fine-grained control over writes

# Introduction: Memory Mapped Files

- + Efficient paging as kernel handles it
- Lesser direct control
- Restricts fine-grained control over writes

# Introduction: Memory Mapped Files

- + Efficient paging as kernel handles it
- Lesser direct control
- Restricts fine-grained control over writes

# Outline

- 1 Introduction
- 2 Architecture**
- 3 Beyond Basics
- 4 Challenges
- 5 Performance Benchmark
- 6 Conclusion
- 7 Questions

# Architecture: Hybrid Log Structured File System (LFS)

- Superblock

- holds references to all other file system data structures
- initializes
  - the inode table
  - the segment controller

- Inodes

- closely resembles inodes from "A Fast File System for UNIX"
- each inode contains 8 direct, 1 indirect, 1 double indirect, and 1 triple indirect blocks
- additionally complicated by the need for an algorithm that allows moving both data blocks and indirect pages belonging to an inode

# Architecture: Hybrid Log Structured File System (LFS)

- Superblock

- holds references to all other file system data structures
- initializes
  - the inode table
  - the segment controller

- Inodes

- closely resembles inodes from "A Fast File System for UNIX"
- each inode contains 8 direct, 1 indirect, 1 double indirect, and 1 triple indirect blocks
- additionally complicated by the need for an algorithm that allows moving both data blocks and indirect pages belonging to an inode

# Architecture: Hybrid Log Structured File System (LFS)

- Superblock

- holds references to all other file system data structures
- initializes
  - the inode table
  - the segment controller

- Inodes

- closely resembles inodes from "A Fast File System for UNIX"
- each inode contains 8 direct, 1 indirect, 1 double indirect, and 1 triple indirect blocks
- additionally complicated by the need for an algorithm that allows moving both data blocks and indirect pages belonging to an inode

# Architecture: Hybrid Log Structured File System (LFS)

- Superblock

- holds references to all other file system data structures
- initializes
  - the inode table
  - the segment controller

- Inodes

- closely resembles inodes from "A Fast File System for UNIX"
- each inode contains 8 direct, 1 indirect, 1 double indirect, and 1 triple indirect blocks
- additionally complicated by the need for an algorithm that allows moving both data blocks and indirect pages belonging to an inode



# Architecture: Hybrid Log Structured File System (LFS)

- Superblock

- holds references to all other file system data structures
- initializes
  - the inode table
  - the segment controller

- Inodes

- closely resembles inodes from "A Fast File System for UNIX"
- each inode contains 8 direct, 1 indirect, 1 double indirect, and 1 triple indirect blocks
- additionally complicated by the need for an algorithm that allows moving both data blocks and indirect pages belonging to an inode

# Architecture: Hybrid Log Structured File System (LFS)

- Superblock

- holds references to all other file system data structures
- initializes
  - the inode table
  - the segment controller

- Inodes

- closely resembles inodes from “A Fast File System for UNIX”
- each inode contains 8 direct, 1 indirect, 1 double indirect, and 1 triple indirect blocks
- additionally complicated by the need for an algorithm that allows moving both data blocks and indirect pages belonging to an inode

# Architecture: Hybrid Log Structured File System (LFS)

- Superblock

- holds references to all other file system data structures
- initializes
  - the inode table
  - the segment controller

- Inodes

- closely resembles inodes from “A Fast File System for UNIX”
- each inode contains 8 direct, 1 indirect, 1 double indirect, and 1 triple indirect blocks
- additionally complicated by the need for an algorithm that allows moving both data blocks and indirect pages belonging to an inode

# Architecture: Hybrid Log Structured File System (LFS)

- Superblock

- holds references to all other file system data structures
- initializes
  - the inode table
  - the segment controller

- Inodes

- closely resembles inodes from “A Fast File System for UNIX”
- each inode contains 8 direct, 1 indirect, 1 double indirect, and 1 triple indirect blocks
- additionally complicated by the need for an algorithm that allows moving both data blocks and indirect pages belonging to an inode

# Architecture: Hybrid Log Structured File System (LFS)

- Superblock

- holds references to all other file system data structures
- initializes
  - the inode table
  - the segment controller

- Inodes

- closely resembles inodes from “A Fast File System for UNIX”
- each inode contains 8 direct, 1 indirect, 1 double indirect, and 1 triple indirect blocks
- additionally complicated by the need for an algorithm that allows moving both data blocks and indirect pages belonging to an inode

# Outline

- 1 Introduction
- 2 Architecture
- 3 Beyond Basics**
- 4 Challenges
- 5 Performance Benchmark
- 6 Conclusion
- 7 Questions

- Segment controller

- heavy lifting data structure
- manages a pool of segments
- tracks segment utilization for efficient GC

- Segment

- each segment maps its data blocks to the inodes that own them
- keeps a write head that allows for constant time chunk allocation
- finding and garbage collecting low utilization chunks can be costly, however

- Experiences with LFS

- + LFS dramatically decreased time to allocate chunks
- + Faster sequential writes due to constant time chunk allocation
- Slower updates due to poor utilization of page cache when doing copy on write for updates
- Introduces costly GC operation

# Beyond Basics

- Segment controller

- heavy lifting data structure
- manages a pool of segments
- tracks segment utilization for efficient GC

- Segment

- each segment maps its data blocks to the inodes that own them
- keeps a write head that allows for constant time chunk allocation
- finding and garbage collecting low utilization chunks can be costly, however

- Experiences with LFS

- + LFS dramatically decreased time to allocate chunks
- + Faster sequential writes due to constant time chunk allocation
- Slower updates due to poor utilization of page cache when doing copy on write for updates
- Introduces costly GC operation



# Beyond Basics

- Segment controller
  - heavy lifting data structure
  - manages a pool of segments
  - tracks segment utilization for efficient GC
- Segment
  - each segment maps its data blocks to the inodes that own them
  - keeps a write head that allows for constant time chunk allocation
  - finding and garbage collecting low utilization chunks can be costly, however
- Experiences with LFS
  - + LFS dramatically decreased time to allocate chunks
  - + Faster sequential writes due to constant time chunk allocation
  - Slower updates due to poor utilization of page cache when doing copy on write for updates
  - Introduces costly GC operation

# Beyond Basics

- Segment controller
  - heavy lifting data structure
  - manages a pool of segments
  - tracks segment utilization for efficient GC
- Segment
  - each segment maps its data blocks to the inodes that own them
  - keeps a write head that allows for constant time chunk allocation
  - finding and garbage collecting low utilization chunks can be costly, however
- Experiences with LFS
  - + LFS dramatically decreased time to allocate chunks
  - + Faster sequential writes due to constant time chunk allocation
  - Slower updates due to poor utilization of page cache when doing copy on write for updates
  - Introduces costly GC operation

# Beyond Basics

- Segment controller
  - heavy lifting data structure
  - manages a pool of segments
  - tracks segment utilization for efficient GC
- Segment
  - each segment maps its data blocks to the inodes that own them
  - keeps a write head that allows for constant time chunk allocation
  - finding and garbage collecting low utilization chunks can be costly, however
- Experiences with LFS
  - + LFS dramatically decreased time to allocate chunks
  - + Faster sequential writes due to constant time chunk allocation
  - Slower updates due to poor utilization of page cache when doing copy on write for updates
  - Introduces costly GC operation

# Beyond Basics

- Segment controller
  - heavy lifting data structure
  - manages a pool of segments
  - tracks segment utilization for efficient GC
- Segment
  - each segment maps its data blocks to the inodes that own them
  - keeps a write head that allows for constant time chunk allocation
  - finding and garbage collecting low utilization chunks can be costly, however
- Experiences with LFS
  - + LFS dramatically decreased time to allocate chunks
  - + Faster sequential writes due to constant time chunk allocation
  - Slower updates due to poor utilization of page cache when doing copy on write for updates
  - Introduces costly GC operation

# Beyond Basics

- Segment controller
  - heavy lifting data structure
  - manages a pool of segments
  - tracks segment utilization for efficient GC
- Segment
  - each segment maps its data blocks to the inodes that own them
  - keeps a write head that allows for constant time chunk allocation
  - finding and garbage collecting low utilization chunks can be costly, however
- Experiences with LFS
  - + LFS dramatically decreased time to allocate chunks
  - + Faster sequential writes due to constant time chunk allocation
  - Slower updates due to poor utilization of page cache when doing copy on write for updates
  - Introduces costly GC operation

# Beyond Basics

- Segment controller
  - heavy lifting data structure
  - manages a pool of segments
  - tracks segment utilization for efficient GC
- Segment
  - each segment maps its data blocks to the inodes that own them
  - keeps a write head that allows for constant time chunk allocation
  - finding and garbage collecting low utilization chunks can be costly, however
- Experiences with LFS
  - + LFS dramatically decreased time to allocate chunks
  - + Faster sequential writes due to constant time chunk allocation
  - Slower updates due to poor utilization of page cache when doing copy on write for updates
  - Introduces costly GC operation

# Beyond Basics

- Segment controller
  - heavy lifting data structure
  - manages a pool of segments
  - tracks segment utilization for efficient GC
- Segment
  - each segment maps its data blocks to the inodes that own them
  - keeps a write head that allows for constant time chunk allocation
  - finding and garbage collecting low utilization chunks can be costly, however
- Experiences with LFS
  - + LFS dramatically decreased time to allocate chunks
  - + Faster sequential writes due to constant time chunk allocation
  - Slower updates due to poor utilization of page cache when doing copy on write for updates
  - Introduces costly GC operation

# Beyond Basics

- Segment controller
  - heavy lifting data structure
  - manages a pool of segments
  - tracks segment utilization for efficient GC
- Segment
  - each segment maps its data blocks to the inodes that own them
  - keeps a write head that allows for constant time chunk allocation
  - finding and garbage collecting low utilization chunks can be costly, however
- Experiences with LFS
  - + LFS dramatically decreased time to allocate chunks
  - + Faster sequential writes due to constant time chunk allocation
  - Slower updates due to poor utilization of page cache when doing copy on write for updates
  - Introduces costly GC operation



# Beyond Basics

- Segment controller
  - heavy lifting data structure
  - manages a pool of segments
  - tracks segment utilization for efficient GC
- Segment
  - each segment maps its data blocks to the inodes that own them
  - keeps a write head that allows for constant time chunk allocation
  - finding and garbage collecting low utilization chunks can be costly, however
- Experiences with LFS
  - + LFS dramatically decreased time to allocate chunks
  - + Faster sequential writes due to constant time chunk allocation
  - Slower updates due to poor utilization of page cache when doing copy on write for updates
  - Introduces costly GC operation

# Beyond Basics

- Segment controller
  - heavy lifting data structure
  - manages a pool of segments
  - tracks segment utilization for efficient GC
- Segment
  - each segment maps its data blocks to the inodes that own them
  - keeps a write head that allows for constant time chunk allocation
  - finding and garbage collecting low utilization chunks can be costly, however
- Experiences with LFS
  - + LFS dramatically decreased time to allocate chunks
  - + Faster sequential writes due to constant time chunk allocation
  - Slower updates due to poor utilization of page cache when doing copy on write for updates
  - Introduces costly GC operation

# Beyond Basics

- Segment controller
  - heavy lifting data structure
  - manages a pool of segments
  - tracks segment utilization for efficient GC
- Segment
  - each segment maps its data blocks to the inodes that own them
  - keeps a write head that allows for constant time chunk allocation
  - finding and garbage collecting low utilization chunks can be costly, however
- Experiences with LFS
  - + LFS dramatically decreased time to allocate chunks
  - + Faster sequential writes due to constant time chunk allocation
  - Slower updates due to poor utilization of page cache when doing copy on write for updates
  - Introduces costly GC operation

# Outline

- 1 Introduction
- 2 Architecture
- 3 Beyond Basics
- 4 Challenges**
- 5 Performance Benchmark
- 6 Conclusion
- 7 Questions

# Challenges

- Memory mapped interface
- Dropped pointer
- Off by one

Language choice: C++ 11

- + proper exception handling
- + proper memory management

# Challenges

- Memory mapped interface
- Dropped pointer
- Off by one

Language choice: C++ 11

- + proper exception handling
- + proper memory management

# Challenges

- Memory mapped interface
- Dropped pointer
- Off by one

Language choice: C++ 11

- + proper exception handling
- + proper memory management

# Challenges

- Memory mapped interface
- Dropped pointer
- Off by one

Language choice: C++ 11

- + proper exception handling
- + proper memory management



# Challenges

- Memory mapped interface
- Dropped pointer
- Off by one

Language choice: C++ 11

- + proper exception handling
- + proper memory management

# Outline

- 1 Introduction
- 2 Architecture
- 3 Beyond Basics
- 4 Challenges
- 5 Performance Benchmark**
- 6 Conclusion
- 7 Questions

# Performance Benchmark

- Synchronous memory mapped
- Asynchronous memory mapped
- Non-memory mapped

# Outline

- 1 Introduction
- 2 Architecture
- 3 Beyond Basics
- 4 Challenges
- 5 Performance Benchmark
- 6 Conclusion**
- 7 Questions

# Conclusion

# Outline

- 1 Introduction
- 2 Architecture
- 3 Beyond Basics
- 4 Challenges
- 5 Performance Benchmark
- 6 Conclusion
- 7 Questions**

# Questions?